



MEDIA INFORMATIC SYSTEMS

IMAGE RECOGNITION TASK

Final progress report

Author

Carlos Sánchez Páez

<http://www.github.com/csp98>

BUDAPEST UNIVERSITY OF TECHNOLOGY AND ECONOMICS

ACADEMIC YEAR 2018-2019

Contents

List of Figures	1
1 Description of the task	2
2 Dataset	2
3 Diary of work	2
4 Preprocessing of the images	3
5 Structure of the neural net	3
6 The final state of the model	5
7 Examples	7
8 Annex: classes contained in the dataset	9
Bibliography	10

List of Figures

1 Loading the CIFAR100 dataset	2
2 Preprocessing of the images.	3
3 Build of the net	4
4 Accuracy and loss of the model (200 epochs)	5
5 Accuracy and loss of the model (139 epochs)	6
6 Bee prediction	7
7 Porcupine prediction	7
8 Bear prediction	8
9 Crocodile prediction	8

1 Description of the task

My task consists on the development of a neural net which will be able to classify different types of RGB images according to the used dataset.

2 Dataset

The train and test images will be taken from *CIFAR100*. It is a dataset which contains 50,000 32x32 color training images, labeled over 100 categories, and 10,000 test images.

The main advantage of this dataset is that it is included in the *keras* dataset. This allows us to load the data by just executing this sentences:

```
from keras.datasets import cifar100
(x_train, y_train), (x_test, y_test) = cifar100.load_data()
```

Figure 1: Loading the CIFAR100 dataset

3 Diary of work

1. **Familiarize with *TensorFlow*.** The first thing I did was reading some *TensorFlow* tutorials. Specifically, I developed a classifier which was able to distinguish between dogs and cats with the help of a tutorial.
2. **Try to adapt the algorithm.** I thought about converting the dogs and cats classifier into a bigger one, so I made the corresponding changes to use it with the *Caltech-101* dataset. The main problem was that the number of images was too low for some categories. That is why I only achieved a 34.5% success rate.
3. **Jump to Caltech-256.** As there were too few images in Caltech-101, I downloaded the Caltech-256 dataset and started to work with it. But I had the opposite problem: as there were too much images, my local PC memory was not enough. To fix this, I created a virtual machine using Google Cloud (30GB of RAM). After 8 hours of training, the accuracy that I achieved was 0.39%. That is because the images for the categories were too similar and I did not apply any preprocessing.
4. **Decide to work with *Caltech-101* and local PC.** As the results obtained in *Caltech-256* were not good, I decided to center myself on *Caltech-101*.
5. **Jump to *Keras*.** In class, we saw a *Keras* example. As I could see, the code was much more readable and easy to understand, so I started to recode the program using *Keras* and *Python 3.6*.
6. **Error with shapes.** I had a lot of problems with the input shapes of my images (I was importing local files directly to the program). I spent here more than a week trying to figure out the error.
7. **Dataset change.** Tired of spending time on the shape error, I decided to use the *CIFAR100* dataset because the loading functions were already implemented.

8. **Parameters calibration.** I modified some parameters to make the model better, such as *batch size* and *preprocessing functions*.

4 Preprocessing of the images

To preprocess the images I use the *ImageDataGenerator* Keras utility. It randomly flips, applies filters, etc. to the images so that the input data is more varied. The parameters used are the following:

```
datagen = ImageDataGenerator(  
    featurewise_center=False,  
    samplewise_center=False,  
    featurewise_std_normalization=False,  
    samplewise_std_normalization=False,  
    zca_whitening=False,  
    rotation_range=0,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True,  
    vertical_flip=True)
```

Figure 2: Preprocessing of the images.

5 Structure of the neural net

- Two convolutional layers with 128 filters using the ELU activation function.
- A pooling layer to reduce the size of the images.
- Two convolutional layers with 256 filters using the ELU activation function.
- A dropout layer to prevent from overfitting.
- Two convolutional layers with 512 filters using the ELU activation function.
- Another dropout layer.
- A flatten layer to convert the result to 1D.
- A fully connected layer.
- Another dropout.
- The output layer with 100 neurons, one for each class. The used activation function is softmax.

The optimizer that the model uses is *RMSProp*, with a learning rate of 0.0001. The batch size is 64.

```

model = Sequential()

model.add(Conv2D(128, (3, 3), padding='same',
                 input_shape=x_train.shape[1:], activation='elu'))
model.add(Conv2D(128, (3, 3), activation='elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(256, (3, 3), padding='same', activation='elu'))
model.add(Conv2D(256, (3, 3), activation='elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(512, (3, 3), padding='same', activation='elu'))
model.add(Conv2D(512, (3, 3), activation='elu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(1024, activation='elu'))
model.add(Dropout(0.5))
model.add(Dense(parameters.NUM_CLASSES, activation='softmax'))

```

Figure 3: Build of the net

6 The final state of the model

First, I trained the model with 200 epochs so that I could find the global maximum of the test accuracy. It took more than 8 hours on my local PC, but the results were quite good:




Figure 4: Accuracy and loss of the model (200 epochs)

We can see that the accuracy is around 60%. Concretely, the maximum accuracy was 68.77%, reached on epoch 139.

With those results, I re-trained the model. It took 3:30 hours and the results were the following:

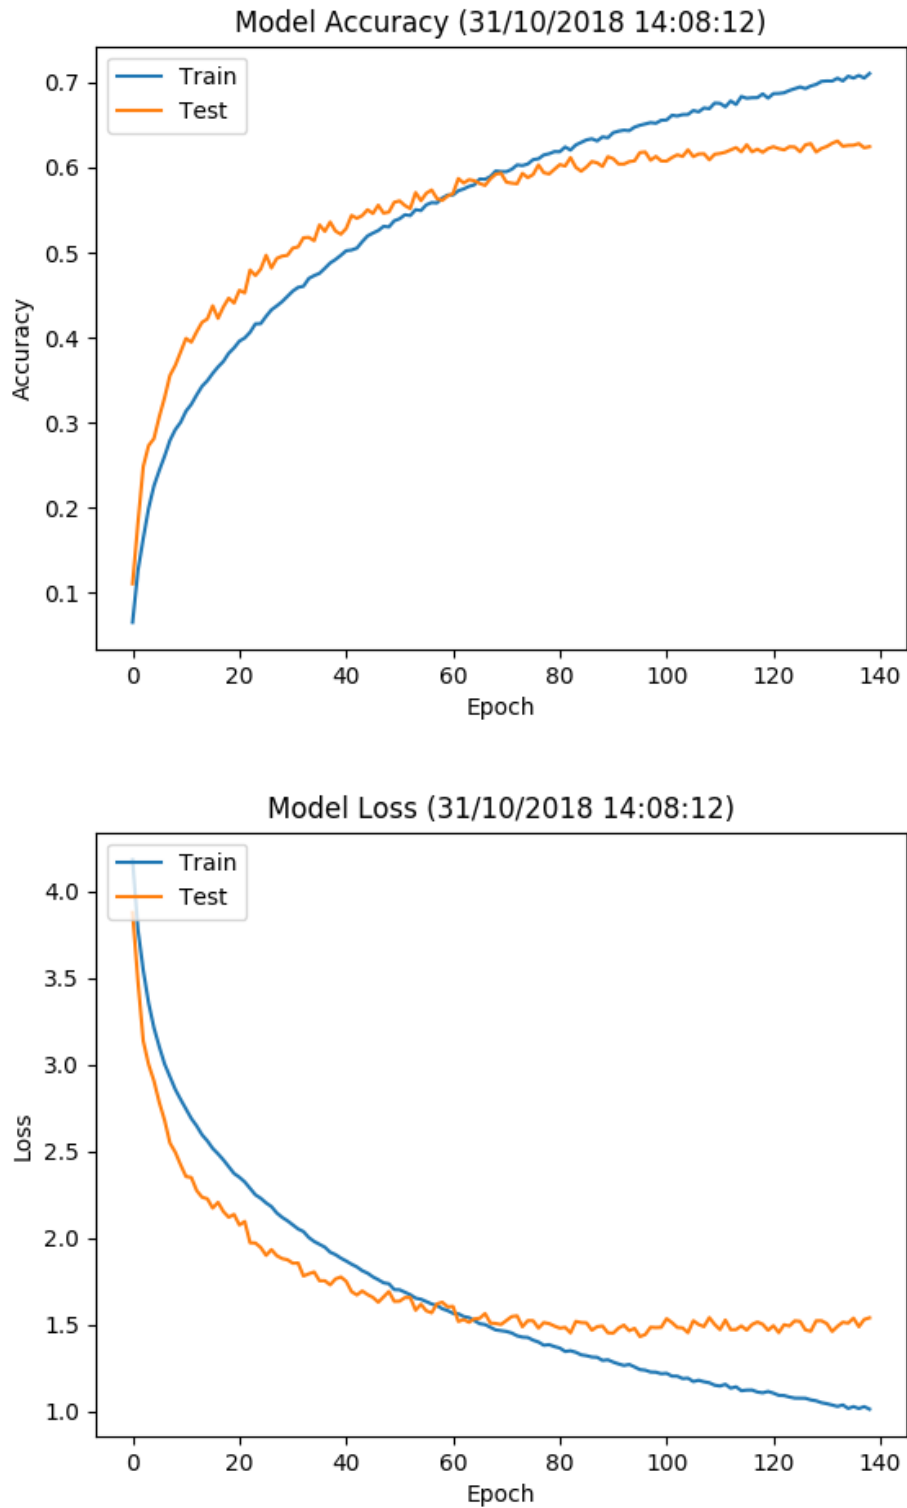


Figure 5: Accuracy and loss of the model (139 epochs)

7 Examples

Let's run the model with different images retrieved from the Internet so that we can see its behaviour.

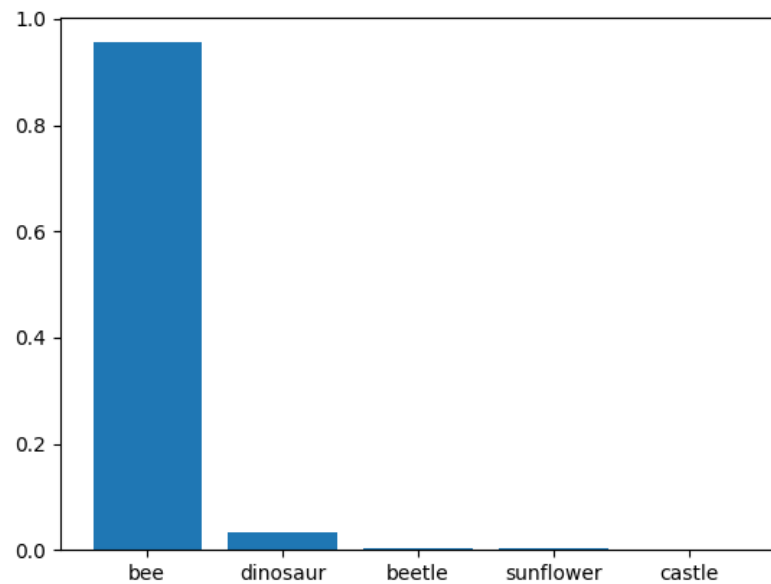


Figure 6: Bee prediction




Figure 7: Porcupine prediction

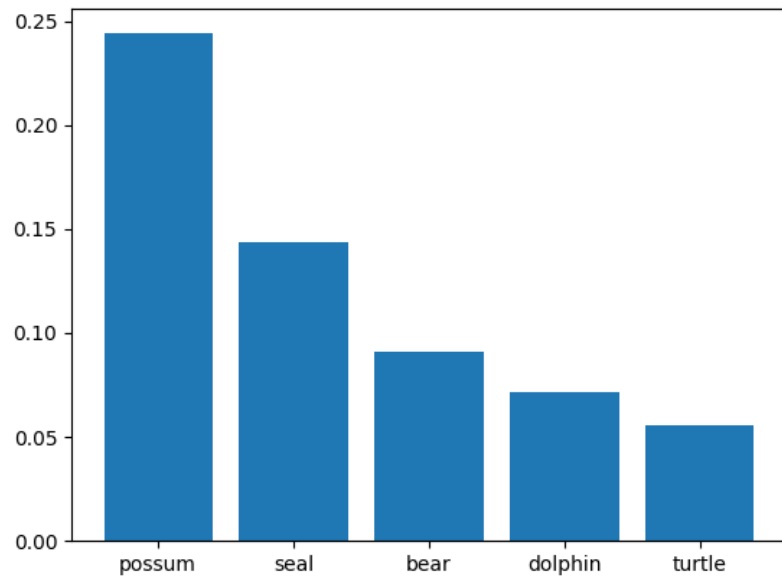



Figure 8: Bear prediction

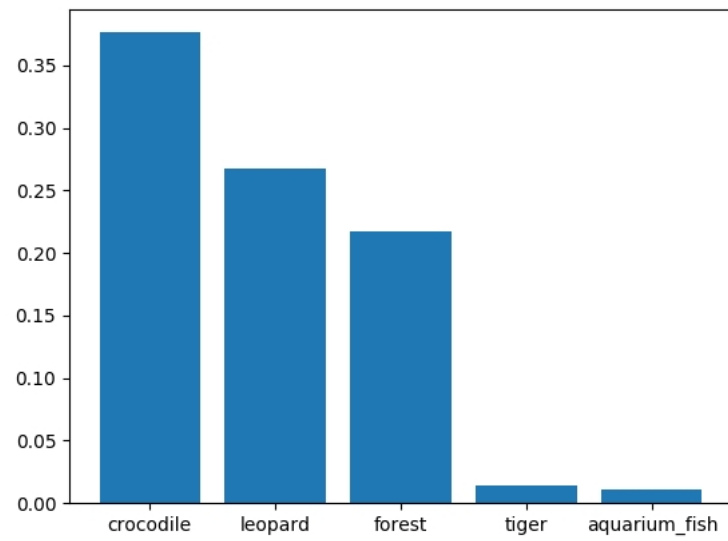


Figure 9: Crocodile prediction

8 Annex: classes contained in the dataset

1. apple	29. cup	57. palm_tree
2. aquarium_fish	30. dinosaur	58. pear
3. baby	31. dolphin	59. pickup_truck
4. bear	32. elephant	60. pine_tree
5. beaver	33. flatfish	61. plain
6. bed	34. forest	62. plate
7. bee	35. fox	63. poppy
8. beetle	36. girl	64. porcupine
9. bicycle	37. hamster	65. possum
10. bottle	38. house	66. rabbit
11. bowl	39. kangaroo	67. raccoon
12. boy	40. keyboard	68. ray
13. bridge	41. lamp	69. road
14. bus	42. lawn_mower	70. rocket
15. butterfly	43. leopard	71. rose
16. camel	44. lion	72. sea
17. can	45. lizard	73. seal
18. castle	46. lobster	74. shark
19. caterpillar	47. man	75. shrew
20. cattle	48. maple_tree	76. skunk
21. chair	49. motorcycle	77. skyscraper
22. chimpanzee	50. mountain	78. snail
23. clock	51. mouse	79. snake
24. cloud	52. mushroom	80. spider
25. cockroach	53. oak_tree	81. squirrel
26. couch	54. orange	82. streetcar
27. crab	55. orchid	83. sunflower
28. crocodile	56. otter	84. sweet_pepper
		85. table
		86. tank

87. telephone	92. trout	97. willow_tree
88. television	93. tulip	98. wolf
89. tiger	94. turtle	99. woman
90. tractor	95. wardrobe	100. worm
91. train	96. whale	

Bibliography

- [1] Keras documentation
<https://keras.io/>
- [2] Moodle
<https://elearning.tmit.bme.hu/login/index.php>
- [3] Tensorflow Tutorial
<https://cv-tricks.com/tensorflow-tutorial/training-convolutional-neural-network-for-image-classification/>
- [4] Some Datasets
<https://www.analyticsvidhya.com/blog/2018/03/comprehensive-collection-deep-learning-datasets/>
- [5] Understanding Convolutions
<https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>
- [6] Deep neural net tutorial
<https://medium.com/@tifa2up/image-classification-using-deep-neural-networks-a-beginner-friendly-approach-using-tensorflow-94b0a090ccd4>
- [7] Caltech-101 dataset
http://www.vision.caltech.edu/Image_Datasets/Caltech101/
- [8] Caltech-256 dataset
http://www.vision.caltech.edu/Image_Datasets/Caltech256/
- [9] TensorFlow tutorial
<https://cv-tricks.com/artificial-intelligence/deep-learning/deep-learning-frameworks/tensorflow/tensorflow-tutorial/>
- [10] Basic classification tutorial with Tensorflow
https://www.tensorflow.org/tutorials/keras/basic_classification
- [11] Keras tutorial
<https://medium.com/@vijayabhaskar96/tutorial-image-classification-with-keras-flow-from-directory-and-generators-95f75ebe5720>

- [12] Another Keras tutorial
<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>
- [13] Doubts resolution
<https://stackoverflow.com/>